

Dialectic Bias in Toxicity Detection of Google’s Perspective API

A study with five parallel corpora

Hanxin Xia | 3417418

hanxin.xia@uni-duesseldorf.de

Abstract

Hate Speech detection has been a heated topic in NLP/NLU since the spread of social media. Despite the existence of various automatic toxicity scoring models, their biases against language varieties that align with minority identities have been discussed by researchers (Sap et al., 2019). In this work, we focus on Google’s Perspective API and analyze its toxicity scores on four synthetic corpora with different dialect features compared to an original master corpus. We discovered that Perspective scored all parallel corpora significantly differently than the master corpus, which further confirms the existence of the biases. Moreover, we discovered the capping phenomenon with the toxicity scores of the gold toxic sentences: Gold non-toxic sentences are more likely to suffer from toxicity score increase, when converted to a dialect, compared to gold toxic sentences.

1 Introduction

With the increasing popularity of social media and the daily widening user bases, the moderation of the contents on online communication platforms have been a crucial part of maintaining the friendliness for the users and building healthy communities. In pursuit of this objective, the platforms have implemented various means including reactive methods such as community guidelines, user report and post-removal of harmful contents and proactive methods either with sensitive words and language pattern rules (Gitari et al., 2015) or machine learning algorithms for automated hate speech detection (Alrehili, 2019).

However current detoxifying approaches are proven to be not perfect. While the reactive methods are commonly subject to deficiency and passivity, the automatic proactive methods often discriminate against minority aligned language varieties other than Standard American English (SAE) (Sap et al., 2019; Zhou et al., 2021).

This research aims to test the performance of Google’s Perspective API on five parallel corpora. Despite being a continuously improving model on toxicity scoring (Google, 2024), the bias against non-standard English is still present in the results. The model scored (state: March 2024) texts from the four English dialects: African American Vernacular English (AAVE), Nigerian and Indian dialects and Singaporean English (Singlish). significantly differently than the original texts of standard English. Thus, we posit that there is a continued need for optimization in the hate speech detection of online contents for better mitigating biases that may infringe upon the online representation of individuals from marginalized communities.

2 Background

The process of detoxification and hate speech detection is aimed to create healthy and welcoming online communities. Yet the biases in these processes are leading to not only unfair penalties for minority individuals (Davidson et al., 2019) but also potential complete exclusion of their voices in online spaces (Blodgett and O’Connor, 2017) simply because of the usage of non-standard language.

The underlying reasons behind such biases can be traced back to multiple factors. As Xia et al. (2020) and Sap et al. (2022) point out, first-hand annotators are inclined to classify non-abusive AAE texts as toxic, which may lead to bias already at the corpus creation stage. Secondly, surface markers on the language such as mention of specific identity terms and usage of certain uncommon patterns are also prone to cause the so-called “key word bias” in models (Resende et al., 2024; Schäfer, 2023). The lack of data of English language varieties in the training data may lead to their incapability to capture the real intentions and deep meanings of such speeches. Last but not the least, suboptimal performances of hate speech detection models can also be attributed to the ignorance of linguistics.

tic subtleties engrained in contextual factors such as speaker identity, environment and even special spelling (Davidson et al., 2019).

Several methodologies have been proposed to mitigate these biases. Schäfer (2023) experimented with transformer based multi-class prediction in order to establish a more comprehensive viewpoint for hate speech detection with less success. Moreover, the approach of masking out identity terms exhibits obvious drawbacks despite its merits. By raising subjectivity level when identity terms occur, (Zhao et al., 2022) were able to efficiently neutralize their model’s bias against minority aligned texts. Similar effects can be achieved by including confident dialect prediction (Ball-Burack et al., 2021). By implementing adversarial learning during the model learning process, Xia et al. (2020) and (Okpala et al., 2022) were both able to reduce baseline model’s false positive rate facing African American English like texts.

3 Data

This section provides an overview on the dataset used in this research. The data augmentation and synthesizing technique will be explained. The choice of the synthetic approach will also be substantiated.

3.1 HateXplain

The data the research is based on stems from the publicly available HateXplain¹ dataset. HateXplain is designed for interpretation of hate speech detection results based on certain text features. It includes 20,148 instances from Twitter and Gab posts collected around the year 2020. Besides the class labeling (*hate*, *offensive* or *normal*), targeted groups and rationales (text chunk that directly contributes to the decision) are also included in the annotation scheme. Each post is annotated by three Amazon Mechanical Turk (MTurk) workers. 919 instances, where all three annotators decide for a different class, are excluded from the final compilation, which ensures the credibility of the annotation results (Mathew et al., 2021).

Since the original class label is tri-partite and the dichotomy of toxicity is interested in the research, we combine *hate* and *offensive* together as toxic, and reinterpret *normal* as non-toxic. Based on majority voting among the three annotators, we are

able to retrieve 12,276 toxic and 7,771 non-toxic posts.

The HateXplain dataset with the newly interpreted labels will be referred to as the **original** or **master dataset** in the following texts.

3.2 Synthetic datasets

Acknowledging the restricted capabilities of current NLP systems on different English language varieties, the Multi-VALUE² package (Ziems et al., 2023) provides a complete framework for rule-based English dialect transformation and evaluation. By applying syntactic and morphological transformation rules based on 189 unique features, Multi-VALUE is able to convert SAE to 50 English dialects.

In this experiment, we choose four representative dialects: African American Vernacular English (AAVE), Nigerian English (NigerianD), Indian English (IndianD) and Colloquial Singapore English (Singlish). The selection includes English dominant communities spanning across three major continents with different historic relations with the English language, which provides good representativeness on the general tendency comparing standard English and dialects.

To measure the distinction of the dialect coded texts, the rules applied on each instance are saved as references. Among the 100 most used rules for each dialect conversion, only 13 of them overlap. This indicates the transformation features are relatively specific to each dialect and inter corpora differences among the synthetic data is high. Examples for most used rules for transformation to each dialect can be found in Appendix B.

In the end, all 20,148 texts from the master dataset were successfully converted to dialect.

3.3 Data choice

Most of the research focused on bias in toxicity detection so far are based on two separate sets of text data (Sap et al., 2019; Davidson et al., 2019; Ball-Burack et al., 2021; Zhou et al., 2021). The data could be sourced from the same domain (i.e. the same social media platform) and split into two subsets. One for standard English and one for texts with certain dialect features. Such an approach ensures the authenticity of the data, since they are real interactions of users online. However, there is

¹<https://github.com/hate-alert/HateXplain>

²<https://github.com/SALT-NLP/multi-value/tree/main>

no real correlation between the texts in those two subsets. At the same time, manual splitting the data requires extra labor. The subjectivity of the annotators may also lead to inaccuracy of categorization of different dialects. Moreover, scarcity has always been an issue while dealing with minority aligned language usage (Dash and Ramamoorthy, 2019). Even on social media, where dialect styled writings are more frequent, the data with such features are still scarce and costly to collect (Jørgensen et al., 2015). These issues are the underlying reasons for the inability to scale in the research focused on dialects and low resource languages.

The advantage of Multi-VALUE in this respect is that we can augment a large number of text data with features specific to certain dialects under the investigation of a study. Using data augmentation to combat scarcity issues is a common approach (Bird et al., 2020). While partly sacrificing the authenticity of the data, Multi-VALUE shows strong capability of preserving meaning and ensuring syntactic integrity and naturalness (Ziems et al., 2022, 2023).

Because the augmented data is 1:1 parallel to that in the master dataset, we are able to investigate the behaviors of one instance with or without dialect features individually. Other than that, by applying multiple conversions, larger scale studies on multiple dialects are also possible, without massive investment in time and energy. Finally, while examining the toxicity score changes after conversion, by backtracking the transformation rules that were used for the conversion, it is also possible to identify the dialect specific features that are responsible for the increase or decrease in text toxicity. This would not be possible if only collected data are used, unless every feature in the dialect data is readily available.

4 Methodology

Perspective API is a multilingual machine learning model for identifying abusive comments developed by Google (Jigsaw, 2017). It has been in constant development since its introduction. The data are sourced from online forums and all manually annotated. The quality of the annotations is ensured by 3-10 annotators per sentence. Being trained on such in-house compiled datasets, Perspective API provides a useful metric to detect hate speech texts from HateXplain dataset without causing data contamination.

All sentences in the five parallel corpora with differently styled English were passed directly to Google’s Perspective API for toxicity scoring using Google Colaboratory platform (access date: March 2024). Because of the 1 request per second rate limit of Perspective, we defined a 0.8 second wait time after successfully retrieving the score of the current sentence. The data are split into 4 batches with 5000 each to address the access time restriction of Google Colab.

5 Results

5.1 Reliability of Perspective API’s scores




Figure 1: Toxicity scores of HateXplain texts by Perspective API across five parallel corpora. 1 = absolute toxic; 0 = not toxic. The scores are split into two subsets based on the gold labels included in the HateXplain annotations.

6 Conclusion

7 Future Studies

References

- Ahlan Alrehili. 2019. [Automatic hate speech detection on social media: A brief survey](#). In *2019 IEEE/ACS 16th International Conference on Computer Systems and Applications (AICCSA)*, pages 1–6.
- Ari Ball-Burack, Michelle Seng Ah Lee, Jennifer Cobbe, and Jatinder Singh. 2021. Differential tweetment: Mitigating racial dialect bias in harmful tweet detection. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 116–128.
- Jordan J. Bird, Diego R. Faria, Cristiano Premebida, Anikó Ekárt, and Pedro P. S. Ayrosa. 2020. [Overcoming data scarcity in speaker identification: Dataset augmentation with synthetic mfccs via character-level rnn](#). In *2020 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*, pages 146–151.
- Su Lin Blodgett and Brendan O’Connor. 2017. Racial disparity in natural language processing: A case study of social media african-american english. *arXiv preprint arXiv:1707.00061*.
- Niladri Sekhar Dash and L. Ramamoorthy. 2019. [Corpus and Dialect Study](#), pages 139–153. Springer Singapore, Singapore.
- Thomas Davidson, Debasmita Bhattacharya, and Ingmar Weber. 2019. [Racial bias in hate speech and abusive language detection datasets](#). In *Proceedings of the Third Workshop on Abusive Language Online*, pages 25–35, Florence, Italy. Association for Computational Linguistics.
- Njagi Dennis Gitari, Zhang Zuping, Hanyurwimfura Damien, and Jun Long. 2015. A lexicon-based approach for hate speech detection. *International Journal of Multimedia and Ubiquitous Engineering*, 10(4):215–230.
- Google. 2024. [Perspective api](#). Accessed: 2024-06-19.
- Jigsaw. 2017. [Perspective api: Using machine learning to reduce online harassment](#). Accessed: 2024-06-28.
- Anna Jørgensen, Dirk Hovy, and Anders Søgaard. 2015. Challenges of studying and processing dialects in social media. In *Proceedings of the workshop on noisy user-generated text*, pages 9–18.
- Binny Mathew, Punyajoy Saha, Seid Muhie Yimam, Chris Biemann, Pawan Goyal, and Animesh Mukherjee. 2021. Hatexplain: A benchmark dataset for explainable hate speech detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 14867–14875.
- Ebuka Okpala, Long Cheng, Nicodemus Mbwapo, and Feng Luo. 2022. Aaebert: Debiasing bert-based hate speech detection models via adversarial learning. In *2022 21st IEEE International Conference on Machine Learning and Applications (ICMLA)*, pages 1606–1612. IEEE.
- Guilherme H Resende, Luiz F Nery, Fabrício Benvenuto, Savvas Zannettou, and Flavio Figueiredo. 2024. A comprehensive view of the biases of toxicity and sentiment analysis methods towards utterances

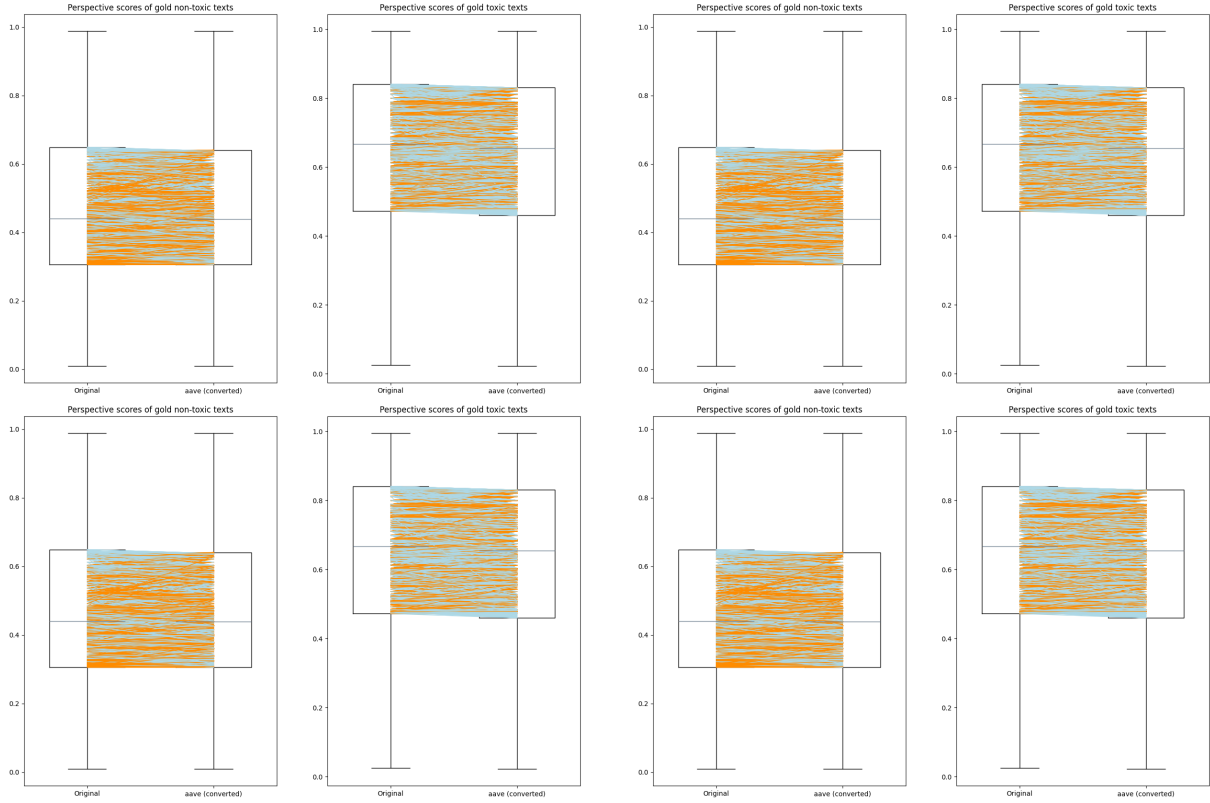


Figure 2

with african american english expressions. *arXiv preprint arXiv:2401.12720*.

Maarten Sap, Dallas Card, Saadia Gabriel, Yejin Choi, and Noah A. Smith. 2019. [The risk of racial bias in hate speech detection](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 1668–1678, Florence, Italy. Association for Computational Linguistics.

Maarten Sap, Swabha Swayamdipta, Laura Vianna, Xuhui Zhou, Yejin Choi, and Noah A. Smith. 2022. [Annotators with attitudes: How annotator beliefs and identities bias toxic language detection](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 5884–5906, Seattle, United States. Association for Computational Linguistics.

Johannes Schäfer. 2023. Bias mitigation for capturing potentially illegal hate speech. *Datenbank-Spektrum*, 23(1):41–51.

Mengzhou Xia, Anjalie Field, and Yulia Tsvetkov. 2020. [Demoting racial bias in hate speech detection](#). In *Proceedings of the Eighth International Workshop on Natural Language Processing for Social Media*, pages 7–14, Online. Association for Computational Linguistics.

Zhixue Zhao, Ziqi Zhang, and Frank Hopfgartner. 2022. [Utilizing subjectivity level to mitigate identity term](#)

[bias in toxic comments classification](#). *Online Social Networks and Media*, 29:100205.

Xuhui Zhou, Maarten Sap, Swabha Swayamdipta, Yejin Choi, and Noah Smith. 2021. [Challenges in automated debiasing for toxic language detection](#). In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 3143–3155, Online. Association for Computational Linguistics.

Caleb Ziems, Jiaao Chen, Camille Harris, Jessica Anderson, and Diyi Yang. 2022. [VALUE: Understanding dialect disparity in NLU](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3701–3720, Dublin, Ireland. Association for Computational Linguistics.

Caleb Ziems, William Held, Jingfeng Yang, Jwala Dhamala, Rahul Gupta, and Diyi Yang. 2023. [Multi-VALUE: A framework for cross-dialectal English NLP](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 744–768, Toronto, Canada. Association for Computational Linguistics.

A Code access

Source code and detailed instructions for conducting analyses are publicly accessible under: <https://github.com/xuanxuanx-98/ToxBias-ParallelCorpora>.

B Rule frequency for dialect conversion

rule name	frequency
yall	4379
demonstrative_for_definite_articles	4197
mass_noun_plurals	4112
regularized_plurals	4083
uninflect	3595
zero_plural	2540
null_relcl	2401
double_modals	2303
drop_copula_be_NP	2294
progressives	2277

Table 1: Top 10 rules used for converting to African American Vernacular English.

rule name	frequency
null_prepositions	8651
mass_noun_plurals	8476
remove_det_definite	7326
progressives	6964
drop_inf_to	5704
yall	2785
regularized_plurals	2660
regularized_past_tense	1636
drop_aux_be_progressive	1524
come_future	1401

Table 2: Top 10 rules used for converting to Nigerian English.

rule name	frequency
progressives	10650
mass_noun_plurals	6902
remove_det_definite	5567
present_perfect_for_past	4820
acomp_focusing_like	4364
definite_for_indefinite_articles	4211
regularized_plurals	4043
null_prepositions	3699
null_referential_pronouns	3675
object_pronoun_drop	3490

Table 3: Top 10 rules used for converting to Indian English.

rule name	frequency
null_prepositions	13860
one_relativizer	11109
zero_plural	8798
plural_to_singular_human	6849
remove_det_definite	6549
mass_noun_plurals	4322
remove_det_indefinite	4078
null_referential_pronouns	3588
object_pronoun_drop	3278
drop_inf_to	3273

Table 4: Top 10 rules used for converting to Colloquial Singapore English.